

Username: Pralay Patoria **Book:** The C++ Standard Library: A Tutorial and Reference, Second Edition. No part of any chapter or book may be reproduced or transmitted in any form by any means without the prior written permission for reprints and excerpts from the publisher of the book or chapter. Redistribution or other use that violates the fair use privilege under U.S. copyright laws (see 17 USC107) or that otherwise violates these Terms of Service is strictly prohibited. Violators will be prosecuted to the full extent of U.S. Federal and Massachusetts laws.

5.3. Numeric Limits

Numeric types in general have platform-dependent limits. The C++ standard library provides these limits in the template `numeric_limits`. These numeric limits replace and supplement the ordinary preprocessor C constants, which are still available for integer types in `<climits>` and `<limits.h>` and for floating-point types in `<cfloat>` and `<float.h>`. The new concept of numeric limits has two advantages: First, it offers more type safety. Second, it enables a programmer to write templates that evaluate these limits.

The numeric limits are discussed in the rest of this section. Note, however, that it is always better to write platform-independent code by using the minimum guaranteed precision of the types. These minimum values are provided in [Table 5.8](#).

¹⁹ Note that “bytes” means an octet with 8 bits. Strictly speaking, it is possible that even a `long int` has one byte with at least 32 bits.

Table 5.8. Minimum Size of Built-In Types

Type	Minimum Size
char	1 byte (8 bits)
short int	2 bytes
int	2 bytes
long int	4 bytes
long long int	8 bytes
float	4 bytes
double	8 bytes
long double	8 bytes

Class `numeric_limits<>`

You usually use templates to implement something once for any type. However, you can also use templates to provide a common interface that is implemented for each type, where it is useful. You can do this by providing specializations of a general template. A typical example of this technique is `numeric_limits`, which works as follows:

- A general template provides the default numeric values for any type:

[Click here to view code image](#)

```
namespace std {
    // general numeric limits as default for any type
    template <typename T>
    class numeric_limits {
    public:
        // by default no specialization for any type T exists
        static constexpr bool is_specialized = false;
        ... // other members are meaningless for the general template
    };
}
```

This general template of the numeric limits says that no numeric limits are available for type `T`. This is done by setting the member `is_specialized` to `false`.

- Specializations of the template define the numeric limits for each numeric type as follows:

```
namespace std {
    // numeric limits for int
    // - implementation defined
    template<> class numeric_limits<int> {
    public:
        // yes, a specialization for numeric limits of int does exist
        static constexpr bool is_specialized = true;

        static constexpr int min() noexcept {
            return -2147483648;
        }
        static constexpr int max() noexcept {
```

```

        return 2147483647;
    }
    static constexpr int digits = 31;
    ...
};
}

```

Here, `is_specialized` is set to `true`, and all other members have the values of the numeric limits for the particular type.

The general `numeric_limits` template and its standard specializations are provided in the header file `<limits>`. The specializations are provided for any fundamental type that can represent numeric values: `bool`, `char`, `signed char`, `unsigned char`, `char16_t`, `char32_t`, `wchar_t`, `short`, `unsigned short`, `int`, `unsigned int`, `long`, `unsigned long`, `long long`, `unsigned long long`, `float`, `double`, and `long double`.²⁰ They can be supplemented easily for user-defined numeric types.

²⁰ The specializations for `char16_t`, `char32_t`, `long long`, and `unsigned long long` are provided since C++11.

Tables 5.9 and 5.10 list all members of the class `numeric_limits<>` and their meanings. Applicable corresponding C constants for these members, defined in `<climits>`, `<limits.h>`, `<float.h>`, and `<float.h>`, are also given.

Table 5.9. Members of Class `numeric_limits<>`, Part 1

Member	Meaning	C Constants
<code>is_specialized</code>	Type has specialization for numeric limits	
<code>is_signed</code>	Type is signed	
<code>is_integer</code>	Type is integer	
<code>is_exact</code>	Calculations produce no rounding errors (true for all integer types)	
<code>is_bounded</code>	The set of values representable is finite (true for all built-in types)	
<code>is_modulo</code>	Adding two positive numbers may wrap to a lesser result	
<code>is_iec559</code>	Conforms to standards IEC 559 and IEEE 754	
<code>min()</code>	Minimum finite value (minimum positive normalized value for floating-point types with denormalization; meaningful if <code>is_bounded</code> <code>is_signed</code>)	<code>INT_MIN</code> , <code>FLT_MIN</code> , <code>CHAR_MIN</code> , ...
<code>max()</code>	Maximum finite value (meaningful if <code>is_bounded</code>)	<code>INT_MAX</code> , <code>FLT_MAX</code> , ...
<code>lowest()</code>	Maximum negative finite value (meaningful if <code>is_bounded</code> ; since C++11)	
<code>digits</code>	Character/integer: number of bits, excluding sign (binary digits)	<code>CHAR_BIT</code>

<code>digits10</code>	Floating point: number of radix digits in the mantissa Number of decimal digits (meaningful if <code>is_bounded</code>)	<code>FLT_MANT_DIG,...</code> <code>FLT_DIG,...</code>
<code>max_digits10</code>	Number of required decimal digits to ensure that values that differ are always differentiated (meaningful for all floating-point types; since C++11)	
<code>radix</code>	Integer: base of the representation (almost always 2) Floating point: base of the exponent representation	<code>FLT_RADIX</code>
<code>min_exponent</code>	Minimum negative integer exponent for base <code>radix</code>	<code>FLT_MIN_EXP,...</code>
<code>max_exponent</code>	Maximum positive integer exponent for base <code>radix</code>	<code>FLT_MAX_EXP,...</code>
<code>min_exponent10</code>	Minimum negative integer exponent for base 10	<code>FLT_MIN_10_EXP,...</code>
<code>max_exponent10</code>	Maximum positive integer exponent for base 10	<code>FLT_MAX_10_EXP,...</code>
<code>epsilon()</code>	Difference of 1 and least value greater than 1	<code>FLT_EPSILON,...</code>
<code>round_style</code>	Rounding style (see page 119)	
<code>round_error()</code>	Measure of the maximum rounding error (according to standard ISO/IEC 10967-1)	
<code>has_infinity</code>	Type has representation for positive infinity	
<code>infinity()</code>	Representation of positive infinity, if available	
<code>has_quiet_NaN</code>	Type has representation for nonsignaling “Not a Number”	
<code>quiet_NaN()</code>	Representation of quiet “Not a Number,” if available	

Table 5.10. Members of Class `numeric_limits<>`, Part 2

Member	Meaning	C Constants
<code>has_signaling_NaN</code>	Type has representation for signaling “Not a Number”	
<code>signaling_NaN()</code>	Representation of signaling “Not a Number,” if available	
<code>has_denorm</code>	Whether type allows denormalized values (variable numbers of exponent bits; see page 119)	
<code>has_denorm_loss</code>	Loss of accuracy is detected as a denormalization loss rather than as an inexact result	
<code>denorm_min()</code>	Minimum positive denormalized value	
<code>traps</code>	Trapping is implemented	
<code>tinyness_before</code>	Tinyness is detected before rounding	

The following is a possible full specialization of the numeric limits for type `float`, which is platform dependent and shows the exact signatures of the members:

[Click here to view code image](#)

```
namespace std {
    template<> class numeric_limits<float> {
    public:
        //yes, a specialization for numeric limits of float does exist
        static constexpr bool is_specialized = true;

        inline constexpr float min() noexcept {
            return 1.17549435E-38F;
        }
        inline constexpr float max() noexcept {
            return 3.40282347E+38F;
        }
        inline constexpr float lowest() noexcept {
            return -3.40282347E+38F;
        }

        static constexpr int digits = 24;
        static constexpr int digits10 = 6;
        static constexpr int max_digits10 = 9;

        static constexpr bool is_signed = true;
```

```

static constexpr bool is_integer = false;
static constexpr bool is_exact = false;
static constexpr bool is_bounded = true;
static constexpr bool is_modulo = false;
static constexpr bool is_iec559 = true;
static constexpr int radix = 2;

inline constexpr float epsilon() noexcept {
    return 1.19209290E-07F;
}

static constexpr float round_style round_style
    = round_to_nearest;
inline constexpr float round_error() noexcept {
    return 0.5F;
}

static constexpr int min_exponent = -125;
static constexpr int max_exponent = +128;
static constexpr int min_exponent10 = -37;
static constexpr int max_exponent10 = +38;

static constexpr bool has_infinity = true;
inline constexpr float infinity() noexcept { return ...; }
static constexpr bool has_quiet_NaN = true;
inline constexpr float quiet_NaN() noexcept { return ...; }
static constexpr bool has_signaling_NaN = true;
inline constexpr float signaling_NaN() noexcept { return ...; }
static constexpr float denorm_style has_denorm = denorm_absent;
static constexpr bool has_denorm_loss = false;
inline constexpr float denorm_min() noexcept { return min(); }

static constexpr bool traps = true;
static constexpr bool tinyness_before = true;
};
}

```

Note that since C++11, all members are declared as **constexpr** ([see Section 3.1.8, page 26](#)). For example, you can use **max()** at places where compile-time expressions are required:

```

static const int ERROR_VALUE = std::numeric_limits<int>::max();
float a[std::numeric_limits<short>::max()];

```

Before C++11, all nonfunction members were constant and static, so their values could be determined at compile time. However, function members were **static** only, so the preceding expressions were not possible. Also note that before C++11, **lowest()** and **max_digits10** were not provided and that empty exception specifications instead of **noexcept** ([see Section 3.1.7, page 24](#)) were used.

The values of **round_style** are shown in [Table 5.11](#). The values of **has_denorm** are shown in [Table 5.12](#). Unfortunately, the member **has_denorm** is not called **denorm_style**. This happened because during the standardization process, there was a late change from a Boolean to an enumerative value. However, you can use the **has_denorm** member as a Boolean value because the standard guarantees that **denorm_absent** is **0**, which is equivalent to **false**, whereas **denorm_present** is **1** and **denorm_indeterminate** is **-1**, both of which are equivalent to **true**. Thus, you can consider **has_denorm** a Boolean indication of whether the type may allow denormalized values.

Table 5.11. Round Style of `numeric_limits<>`

Round Style	Meaning
<code>round_toward_zero</code>	Rounds toward zero
<code>round_to_nearest</code>	Rounds to the nearest representable value
<code>round_toward_infinity</code>	Rounds toward positive infinity
<code>round_toward_neg_infinity</code>	Rounds toward negative infinity
<code>round_indeterminate</code>	Indeterminable

Table 5.12. Denormalization Style of `numeric_limits<>`

Denorm Style	Meaning
denorm_absent	The type does not allow denormalized values
denorm_present	The type allows denormalized values to the nearest representable value
denorm_indeterminate	Indeterminable

Example of Using `numeric_limits`

The following example shows possible uses of some numeric limits, such as the maximum values for certain types and determining whether `char` is signed:

[Click here to view code image](#)

```
// util/limits1.cpp

#include <iostream>
#include <limits>
#include <string>
using namespace std;

int main()
{
    // use textual representation for bool
    cout << boolalpha;

    // print maximum of integral types
    cout << "max(short): " << numeric_limits<short>::max() << endl;
    cout << "max(int): " << numeric_limits<int>::max() << endl;
    cout << "max(long): " << numeric_limits<long>::max() << endl;
    cout << endl;

    // print maximum of floating-point types
    cout << "max(float): " << numeric_limits<float>::max() << endl;
    cout << "max(double): " << numeric_limits<double>::max() << endl;
    cout << "max(long double): " << numeric_limits<long double>::max() << endl;
    cout << endl;

    // print whether char is signed
    cout << "is_signed(char): " << numeric_limits<char>::is_signed << endl;
    cout << endl;

    // print whether numeric limits for type string exist
    cout << "is_specialized(string): " << numeric_limits<string>::is_specialized << endl;
}
```

The output of this program is platform dependent. Here is a possible output of the program:

```
max(short): 32767
max(int): 2147483647
max(long): 2147483647

max(float): 3.40282e+38
max(double): 1.79769e+308
max(long double): 1.79769e+308

is_signed(char): false

is_specialized(string): false
```

The last line shows that no numeric limits are defined for the type `string`. This makes sense because strings are not numeric values. However, this example shows that you can query for any arbitrary type whether or not it has numeric limits defined.