

# Operators



The “SQL equivalents” column in the reference tables in this chapter do not necessarily correspond to what an `IQueryable` implementation such as LINQ to SQL will produce. Rather, it indicates what you’d typically use to do the same job if you were writing the SQL query yourself. Where there is no simple translation, the column is left blank. Where there is no translation at all, the column reads “Exception thrown”.

`Enumerable` implementation code, when shown, excludes checking for null arguments and indexing predicates.

With each of the filtering methods, you always end up with either the same number

With each of the filtering methods, you always end up with either the same number or fewer elements than you started with. You can never get more! The elements are also identical when they come out; they are not transformed in any way.

# Where

Argument	Type
----------	------

Source sequence	<code>IEnumerable&lt;TSource&gt;</code>
-----------------	---

Predicate	<code>TSource =&gt; bool</code> or <code>(TSource, int) =&gt; bool<sup>a</sup></code>
-----------	---

<sup>a</sup> Prohibited with LINQ to SQL and Entity Framework.

## Query syntax

where *bool-expression*

Enumerable.Where(Whereable WhereableImplementation)

# Enumerable.Where implementation

The internal implementation of Enumerable.Where, null checking aside, is functionally equivalent to the following:

```
public static IEnumerable<TSource> Where<TSource>
    (this IEnumerable<TSource> source, Func<TSource, bool> predicate)
{
    foreach (TSource element in source)
        if (predicate(element))
            yield return element;
}
```

## Overview

Where returns the elements from the input sequence that satisfy the given predicate.

For instance:

```
string[] names = { "Tom", "Dick", "Harry", "Mary", "Jay" };  
IEnumerable<string> query = names.Where (name => name.EndsWith ("y"));
```

```
// Result: { "Harry", "Mary", "Jay" }
```

---

372 | Chapter 9: LINQ Operators

## In query syntax:

```
IEnumerable<string> query = from n in names  
                             where n.EndsWith ("y")  
                             select n;
```

A **where** clause can appear more than once in a query and be interspersed with **let** clauses:

```
from n in names
```

```
from n in names
where n.Length > 3
let u = n.ToUpper()
where u.EndsWith ("Y")
select u;
```

```
// Result: { "HARRY", "MARY" }
```

Standard C# scoping rules apply to such queries. In other words, you cannot refer to a variable prior to declaring it with an range variable or a let clause.

## Indexed filtering

Where's predicate optionally accepts a second argument, of type int. This is fed with the position of each element within the input sequence, allowing the predicate to use this information in its filtering decision. For example, the following skips every

the position of each element within the input sequence, allowing the predicate to use this information in its filtering decision. For example, the following skips every second element:

```
IEnumerable<string> query = names.Where ((n, i) => i % 2 == 0);  
  
// Result: { "Tom", "Harry", "Jay" }
```

An exception is thrown if you use indexed filtering in LINQ to SQL or EF.

## **SQL LIKE comparisons in LINQ to SQL and EF**

The following methods on string translate to SQL's LIKE operator:

**Contains, StartsWith, EndsWith**

For instance, `c.Name.Contains ("abc")` translates to `customer.Name LIKE '%abc%'` (or more accurately, a parameterized version of this). `Contains` lets you compare only a against a locally evaluated expression; to compare against another column, you must use the `SqlMethods.Like` method:

against a locally evaluated expression; to compare against another column, you must use the `SqlMethods.Like` method:

```
... where SqlMethods.Like (c.Description, "%" + c.Name + "%")
```

`SqlMethods.Like` also lets you perform more complex comparisons (e.g., `LIKE 'abc%def%'`).

## < and > string comparisons in LINQ to SQL and EF

You can perform *order* comparison on strings with `string's CompareTo` method; this maps to SQL's < and > operators:

```
dataContext.Purchases.Where (p => p.Description.CompareTo ("C") < 0)
```

## Operators

### **WHERE x IN (...; ...; ...) in LINQ to SQL and EF**

With LINQ to SQL and EF, you can apply the Contains operator to a local collection within a filter predicate. For instance:

```
string[] chosenNames = { "Tom", "Jay" };
```

```
from c in DataContext.Customers
where chosenNames.Contains (c.Name)
...
```

This maps to SQL's IN operator—in other words:



This maps to SQL's `IN` operator—in other words:

```
WHERE customer.Name IN ("Tom", "Jay")
```

If the local collection is an array of entities or nonscalar types, LINQ to SQL or EF may instead emit an `EXISTS` clause.

## Take and Skip

Argument	Type
Source sequence	<code>IEnumerable&lt;TSource&gt;</code>
Number of elements to take or skip	<code>int</code>

`Take` emits the first  $n$  elements and discards the rest; `Skip` discards the first  $n$  elements and emits the rest. The two methods are useful together when implementing a web page allowing a user to navigate through a large set of matching records. For instance, suppose a user searches a book database for the term “mercury,” and there are 100 matches. The following returns the first 20:

```
IQueryable<Book> query = dbContext.Books  
    .Where (b => b.Title.Contains ("mercury"))  
    .OrderBy (b => b.Title)  
    .Take (20);
```

The next query returns books 21 to 40:

```
IQueryable<Book> query = dbContext.Books  
    .Where (b => b.Title.Contains ("mercury"))  
    .OrderBy (b => b.Title)  
    .Skip (20).Take (20);
```

LINQ to SQL and EF translate Take and Skip to the ROW\_NUMBER function in SQL Server 2005, or a TOP *n* subquery in earlier versions of SQL Server.

## TakeWhile and SkipWhile

# TakeWhile and SkipWhile

Argument	Type
Source sequence	IEnumerable<TSource>
Predicate	TSource => bool or (TSource, int) => bool

---

## 374 | Chapter 9: LINQ Operators

---

**TakeWhile** enumerates the input sequence, emitting each item, until the given predicate is false. It then ignores the remaining elements:

```
int[] numbers = { 3, 5, 2, 234, 4, 1 };  
var takeWhileSmall = numbers.TakeWhile (n => n < 100);    // { 3, 5, 2 }
```

**SkipWhile** enumerates the input sequence, ignoring each item until the given predicate is false. It then emits the remaining elements:

```
int[] numbers = { 3, 5, 2, 234, 4, 1 };  
var skipWhileSmall = numbers.SkipWhile (n => n < 100);    // { 234, 4, 1 }
```

**TakeWhile** and **SkipWhile** have no translation to SQL and cause a runtime error if

`TakeWhile` and `SkipWhile` have no translation to SQL and cause a runtime error if used in a LINQ-to-db query.

# Distinct

`Distinct` returns the input sequence, stripped of duplicates. Only the default equality comparer can be used for equality comparison. The following returns distinct letters in a string:

```
char[] distinctletters = "HelloWorld".Distinct().ToArray();  
string s = new string (distinctletters);           // HelloWorld
```

We can call LINQ methods directly on a string, because string implements `IEnumerable<char>`.

# Projecting

`Enumerable<TSource>→IEnumerable<TResult>`

Method	Description	SQL equivalents
Select	Transforms each input element with the given lambda expression	SELECT
SelectMany	Transforms each input element, and then flattens and concatenates the resultant subsequences	INNER JOIN, LEFT OUTER JOIN,  CROSS JOIN



When querying a database, **Select** and **SelectMany** are the most versatile joining constructs; for local queries, **Join** and **Group Join** are the most *efficient* joining constructs.

# Select

Argument	Type
Source sequence	<code>IEnumerable&lt;TSource&gt;</code>
Result selector	<code>TSource =&gt; TResult</code> or <code>(TSource, int) =&gt; TResult</code> <sup>a</sup>

Result selector      TSource => TResult<sup>1</sup> or (TSource, int) => TResult<sup>2</sup>

---

a Prohibited with LINQ to SQL and Entity Framework.

---

Projecting | 375

A black rectangular box with the text "LINQ Operators" in white, sans-serif font, centered within the box.

LINQ Operators

# Query syntax

Queryable

*select projection-expression*

## Enumerable implementation

```
public static IEnumerable<TResult> Select<TSource, TResult>
    (this IEnumerable<TSource> source, Func<TSource, TResult> selector)
{
    foreach (TSource element in source)
        yield return selector(element);
}
```

## Overview

With Select, you always get the same number of elements that you started with. Each element, however, can be transformed in any manner by the lambda function.

The following selects the names of all fonts installed on the computer (from

The following selects the names of all fonts installed on the computer (from `System.Drawing`):

```
IEnumerable<string> query = from f in FontFamily.Families
                             select f.Name;
```

```
foreach (string name in query) Console.WriteLine (name);
```

In this example, the select clause converts a `FontFamily` object to its name. Here's the lambda equivalent:

```
IEnumerable<string> query = FontFamily.Families.Select (f => f.Name);
```

Select statements are often used to project into anonymous types:

```
var query =
    from f in FontFamily.Families
    select new { f.Name, LineSpacing = f.GetLineSpacing (FontStyle.Bold) };
```

A projection with no transformation is sometimes used with query syntax, in order to satisfy the requirement that the query end in a select or group clause. The following selects fonts supporting strikethrough:



to satisfy the requirement that the query end in a select or group clause. The following selects fonts supporting **strikeout**:

```
IEnumerable<FontFamily> query =  
    from f in FontFamily.Families  
    where f.IsStyleAvailable (FontStyle.Strikeout)  
    select f;
```

```
foreach (FontFamily ff in query) Console.WriteLine (ff.Name);
```

In such cases, the compiler omits the projection when translating to fluent syntax.

## Indexed projection

The selector expression can optionally accept an integer argument, which acts as an indexer, providing the expression with the position of each input in the input sequence. This works only with local queries:

```
string[] names = { "Tom", "Dick", "Harry", "Mary", "Jay" };
```

```
IEnumerable<string> query = names
    .Select ((s,i) => i + "=" + s);

//
{ "0=Tom", "1=Dick",
  ... }
```

## Select subqueries and object hierarchies

You can nest a subquery in a select clause to build an object hierarchy. The following example returns a collection describing each directory under *D:\source*, with a subcollection of files under each directory:

lowing example returns a collection describing each directory under *D:\source*, with a subcollection of files under each directory:

```
DirectoryInfo[] dirs = new DirectoryInfo(@"d:\source").GetDirectories();
```

```
var query =
```

```
from d in dirs
```

```
where (d.Attributes & FileAttributes.System) == 0
```

```
select new
```

```
{
```

```
    DirectoryName = d.FullName,
```

```
    Created = d.CreationTime,
```

```
    Files = from f in d.GetFiles()
```

```
        where (f.Attributes & FileAttributes.Hidden) == 0
```

```
        select new { FileName = f.Name, f.Length, }
```

```
};
```

```
foreach (var dirFiles in query)
```

```
{
```

```
    Console.WriteLine("Directory: " + dirFiles.DirectoryName);
```

```
{  
    Console.WriteLine ("Directory: " + dirFiles.DirectoryName);  
    foreach (var file in dirFiles.Files)  
        Console.WriteLine (" " + file.FileName + "len: " + file.Length);  
}
```

The inner portion of this query can be called a *correlated subquery*. A subquery is correlated if it references an object in the outer query—in this case, it references *d*, the directory being enumerated.



A subquery inside a *Select* allows you to map one object hierarchy to another, or map a relational object model to a hierarchical object model.

With local queries, a subquery within a *Select* causes double-deferred execution. In our example, the files don't get filtered or projected until the inner *foreach* statement enumerates.

### **Subqueries and joins in LINQ to SQL and EF**

Subquery projections work well in LINQ to SQL and EF and can be used to do the

Subquery projections work well in LINQ to SQL and EF and can be used to do the work of SQL-style joins. Here's how we retrieve each customer's name along with their high-value purchases:

---

Projecting | 377

## LINQ Operators

```
var query =  
    from c in dataContext.Customers  
    select new {
```

```
from c in dataContext.Customers  
select new {
```

```
    c.Name,
```

```
    Purchases = from p in dataContext.Purchases
```

```
        where p.CustomerID == c.ID && p.Price > 1000
```

```
        select new { p.Description, p.Price }  
};
```

```
foreach (var namePurchases in query)
```


```
{
```

```
    Console.WriteLine ("Customer: " + namePurchases.Name);
```

```
    foreach (var purchasedDetail in namePurchases.Purchases)
```

```
        Console.WriteLine ("    - $$$: " + purchasedDetail.Price);
```

```
}
```



This style of query is ideally suited to interpreted queries. The outer query and subquery are processed as a unit, avoiding unnecessary round-tripping. With local queries, however, it's inefficient because every combination of outer and inner elements must be enumerated to get the few matching combinations. A better choice for local queries is `Join` or `GroupJoin`, described in the following sections.

This query matches up objects from two disparate collections, and it can be thought of as a “Join.” The difference between this and a conventional database join (or subquery) is that we're not flattening the output into a single two-dimensional result set. We're mapping the relational data to hierarchical data, rather than to flat data.

Here's the same query simplified by using the `Purchases` association property on the `Customer` entity:

```
from c in dataContext.Customers
select new
{
    c.Name,
```

```
{  
    c.Name,  
    Purchases = from p in c.Purchases    // Purchases is EntitySet<Purchase>  
                where p.Price > 1000  
                select new { p.Description, p.Price }  
};
```

Both queries are analogous to a left outer join in SQL in the sense that we get all customers in the outer enumeration, regardless of whether they have any purchases. To emulate an inner join—where customers without high-value purchases are excluded—we would need to add a filter condition on the purchases collection:

```
from c in DataContext.Customers  
where c.Purchases.Any (p => p.Price > 1000)  
select new {  
    c.Name,  
    Purchases = from p in c.Purchases  
                where p.Price > 1000  
                select new { p.Description, p.Price }  
};
```



This is slightly untidy, however, in that we've written the same predicate (`Price > 1000`) twice. We can avoid this duplication with a `let` clause:

```
from c in dataContext.Customers
let highValueP = from p in c.Purchases
                  where p.Price > 1000
                  select new { p.Description, p.Price }
where highValueP.Any()
select new { c.Name, Purchases = highValueP };
```

This style of query is flexible. By changing `Any` to `Count`, for instance, we can modify the query to retrieve only customers with at least two high-value purchases:

```
...
where highValueP.Count() >= 2
select new { c.Name, Purchases = highValueP };
```

### Projecting into concrete types

Projecting into anonymous types is useful in obtaining intermediate results, but not so useful if you want to send a result set back to a client, for instance, because

so useful if you want to send a result set back to a client, for instance, because anonymous types can exist only as local variables within a method. An alternative is to use concrete types for projections, such as `DataSets` or custom business entity classes. A custom business entity is simply a class that you write with some properties, similar to a LINQ to SQL `[Table]` annotated class or an EF Entity, but designed to hide lower-level (database-related) details. You might exclude foreign key fields from business entity classes, for instance. Assuming we wrote custom entity classes called `CustomerEntity` and `PurchaseEntity`, here's how we could project into them:

```
Queryable<CustomerEntity> query =  
    from c in DataContext.Customers  
    select new CustomerEntity  
    {  
        Name = c.Name,  
        Purchases =  
            (from p in c.Purchases  
             where p.Price > 1000  
             select new PurchaseEntity {  
                 Description = p.Description,  
                 Value = p.Price  
             })  
    }.ToList()  
};
```

// Force query execution, converting output to a more convenient list:

```
// Force query execution, converting output to a more convenient List:  
List<CustomerEntity> result = query.ToList();
```

Notice that so far, we've not had to use a `Join` or `SelectMany` statement. This is because we're maintaining the hierarchical shape of the data, as illustrated in Figure 9-2. With LINQ, you can often avoid the traditional SQL approach of flattening tables into a two-dimensional result set.




Figure 9-2. Projecting an object hierarchy

# SelectMany

Argument	Type
----------	------

Source sequence	IEnumerable<TSource>
-----------------	----------------------

Result selector	TSource => IEnumerable<TResult>
-----------------	---------------------------------

or (TSource, int) => IEnumerable<TResult><sup>a</sup>

a Prohibited with LINQ to SQL

## Query syntax

*from identifier1 in enumerable-expression1*  
*from identifier2 in enumerable-expression2*  
...

## Enumerable implementation

# Enumerable Implementation

```
public static IEnumerable<TResult> SelectMany<TSource, TResult>
    (IEnumerable<TSource> source,
     Func<TSource, IEnumerable<TResult>> selector)
{
    foreach (TSource element in source)
        foreach (TResult subElement in selector (element))
            yield return subElement;
}
```

## Overview

SelectMany concatenates subsequences into a single flat output sequence.

Recall that for each input element, Select yields exactly one output element. In contrast, SelectMany yields  $0..n$  output elements. The  $0..n$  elements come from a subsequence or child sequence that the lambda expression must emit.

SelectMany can be used to expand child sequences, flatten nested collections, and join two collections into a flat output sequence. Using the conveyor belt analogy,

`SelectMany` can be used to expand child sequences, flatten nested collections, and join two collections into a flat output sequence. Using the conveyor belt analogy, `SelectMany` funnels fresh material onto a conveyor belt. With `SelectMany`, each input

---

## 380 | Chapter 9: LINQ Operators

element is the *trigger* for the introduction of fresh material. The fresh material is emitted by the selector lambda expression and must be a sequence. In other words, the lambda expression must emit a *child sequence* per input *element*. The final result is a concatenation of the child sequences emitted for each input element.

Starting with a simple example, suppose we have an array of names as follows:

```
string[] fullNames = { "Anne Williams", "John Fred Smith", "Sue Green" };
```

which we wish to convert to a single flat collection of words—in other words:

```
"Anne", "Williams", "John", "Fred", "Smith", "Sue", "Green"
```

`SelectMany` is ideal for this task, because we're mapping each input element to a variable number of output elements. All we must do is come up with a selector expression that converts each input element to a child sequence. `string.Split` does the job nicely: it takes a string and splits it into words, emitting the result as an array.

expression that converts each input element to a child sequence. `string.Split` does the job nicely: it takes a string and splits it into words, emitting the result as an array:

```
string testInputElement = "Anne Williams";  
string[] childSequence = testInputElement.Split();
```

// childSequence is { "Anne", "Williams" };

So, here's our `SelectMany` query and the result:

```
IEnumerable<string> query = fullNames.SelectMany (name => name.Split());  
  
foreach (string name in query)
```

```
    Console.Write (name + "|"); // Anne|Williams|John|Fred|Smith|Sue|Green|
```



If you replace `SelectMany` with `Select`, you get the same results in hierarchical form. The following emits a sequence of string *arrays*, requiring nested `foreach` statements to enumerate:

```
IEnumerable<string[]> query =  
    fullNames.Select (name => name.Split());
```



```
foreach (string[] stringArray in query)
    foreach (string name in stringArray)
        Console.WriteLine (name + "/" );
```

The benefit of `SelectMany` is that it yields a single *flat* result sequence.

`SelectMany` is supported in query syntax and is invoked by having an *additional generator*—in other words, an extra `from` clause in the query. The `from` keyword has two meanings in query syntax. At the start of a query, it introduces the original range variable and input sequence. *Anywhere else* in the query, it translates to `SelectMany`. Here's our query in query syntax:

```
IEnumerable<string> query =
    from fullName in fullNames
    from name in fullName.Split()
```

```
from name in FullName.Split()  
select name;  
  
// Translates to SelectMany
```

Note that the additional generator introduces a new query variable—in this case, `name`. The new query variable becomes the range variable from then on, and the old range variable is demoted to an *outer range variable*.

## Outer range variables

In the preceding example, `fullName` becomes an outer range variable after `SelectMany`. Outer range variables remain in scope until the query either ends or reaches an `into` clause. The extended scope of these variables is *the* killer scenario for query syntax over fluent syntax.

To illustrate, we can take the preceding query and include `fullName` in the final projection:

```
IEnumerable<string> query =  
    from fullName in fullNames           // fullName = outer variable  
    from name in fullName.Split()       // name = range variable  
    select name + " came from " + fullName;
```

Anne came from Anne Williams

Anne came from Anne Williams  
Williams came from Anne Williams  
John came from John Fred Smith  
...

Behind the scenes, the compiler must pull some tricks to resolve outer references. A good way to appreciate this is to try writing the same query in fluent syntax. It's tricky! It gets harder still if you insert a `where` or `orderby` clause before projecting:

```
from fullName in fullNames
from name in fullName.Split()
orderby fullName, name
select name + " came from " + fullName;
```

The problem is that `SelectMany` emits a flat sequence of child elements—in our case, a flat collection of words. The original outer element from which it came (`fullName`) is lost. The solution is to “carry” the outer element with each child, in a

a list collection of words. The original outer element from which it came (`fullName`) is lost. The solution is to “carry” the outer element with each child, in a temporary anonymous type:

```
from fullName in fullNames
from x in fullName.Split().Select (name => new { name, fullName } )
orderby x.fullName, x.name
select x.name + " came from " + x.fullName;
```

The only change here is that we’re wrapping each child element (`name`) in an anonymous type that also contains its `fullName`. This is similar to how a `let` clause is resolved. Here’s the final conversion to fluent syntax:

```
IEnumerable<string> query = fullNames
    .SelectMany (fName => fName.Split()
        .Select (name => new { name, fName } ))
    .OrderBy (x => x.fName)
    .ThenBy (x => x.name)
    .Select (x => x.name + " came from " + x.fName);
```

# Thinking in query syntax

As we just demonstrated, there are good reasons to use query syntax if you need the outer range variable. In such cases, it helps not only to use query syntax, but also to think directly in its terms.

There are two basic patterns when writing additional generators. The first is *expanding and flattening subsequences*. To do this, you call a property or method on an existing query variable in your additional generator. We did this in the previous example:

```
from fullName in fullNames
from name in fullName.Split()
```

Here, we've expanded from enumerating full names to enumerating words. An analogous LINQ-to-db query is when you expand child association properties. The following query lists all customers along with their purchases:

```
IEnumerable<string> query = from c in DataContext.Customers  
                             from p in c.Purchases  
                             select c.Name + " bought a " + p.Description;
```

Tom bought a Bike  
Tom bought a Holiday  
Dick bought a Phone  
Harry bought a Car  
...

Here, we've expanded each customer into a subsequence of purchases.

The second pattern is performing a *cross product* or *cross join*—where every element of one sequence is matched with every element of another. To do this, introduce a generator whose selector expression returns a sequence unrelated to a range variable:

```
int[] numbers = { 1, 2, 3 };  
  
string[] letters = { "a", "b" };
```

```
IEnumerable<string> query = from n in numbers  
                             from l in letters  
                             select n.ToString() + l;
```

```
RESULT: { "1a", "1b", "2a", "2b", "3a", "3b" }
```

This style of query is the basis of *SelectMany*-style *joins*.

## Joining with *SelectMany*

You can use *SelectMany* to join two sequences, simply by filtering the results of a cross product. For instance, suppose we wanted to match players for a game. We could start as follows:



cross product. For instance, suppose we wanted to match players for a game. We could start as follows:

```
string[] players = { "Tom", "Jay", "Mary" };
```

```
IEnumerable<string> query = from name1 in players  
                             from name2 in players  
                             select name1 + " vs " + name2;
```

```
RESULT: { "Tom vs Tom", "Tom vs Jay", "Tom vs Mary",  
          "Jay vs Tom", "Jay vs Jay", "Jay vs Mary",  
          "Mary vs Tom", "Mary vs Jay", "Mary vs Mary" }
```

The query reads: "For every player, reiterate every player, selecting player 1 vs player 2". Although we got what we asked for (a cross join), the results are not useful until we add a filter:

```
IEnumerable<string> query = from name1 in players  
                             from name2 in players  
                             where name1.CompareTo(name2) < 0  
                             orderby name1, name2  
                             select name1 + " vs " + name2;
```

```
RESULT: { "Jay vs Mary", "Jay vs Tom", "Mary vs Tom" }
```

The filter predicate constitutes the *join condition*. Our query can be called a *non-equijoin*, because the join condition doesn't use an equality operator.

We'll demonstrate the remaining types of joins with LINQ to SQL (they'll also work with EF except where we explicitly use a foreign key field).

## SelectMany in LINQ to SQL and EF

SelectMany in LINQ to SQL and EF can perform cross joins, non-equijoin, inner joins, and left outer joins. You can use SelectMany with both predefined associations and ad hoc relationships—just as with Select. The difference is that SelectMany returns a flat rather than a hierarchical result set.

A LINQ-to-db cross join is written just as in the preceding section. The following query matches every customer to every purchase (a cross join):

```
var query = from c in dataContext.Customers
             from p in dataContext.Purchases
             select c.Name + " might have bought a " + p.Description;
```

```
select c.Name + " might have bought a " + p.Description;
```

More typically, though, you'd want to match customers to their own purchases only. You achieve this by adding a `where` clause with a joining predicate. This results in a standard SQL-style equi-join:

```
var query = from c in dataContext.Customers
             from p in dataContext.Purchases
             where c.ID == p.CustomerID
             select c.Name + " bought a " + p.Description;
```



This translates well to SQL. In the next section, we'll see how it extends to support outer joins. Reformulating such queries with LINQ's `Join` operator actually makes them *less* extensible—LINQ is opposite to SQL in this sense.

If you have association properties for relationships in your entities, you can express the same query by expanding the subcollection instead of filtering the cross product:

```
from c in DataContext.Customers
from p in c.Purchases
select new { c.Name, p.Description };
```



Entity Framework doesn't expose foreign keys in the entities, so for recognized relationships you *must* use its association properties rather than joining manually as we did previously.

The advantage is that we've eliminated the joining predicate. We've gone from filtering a cross product to expanding and flattening. Both queries, however, will result in the same SQL.

You can add `where` clauses to such a query for additional filtering. For instance, if we wanted only customers whose names started with "T", we could filter as follows:

```
from c in DataContext.Customers
```

```
from c in DataContext.Customers
where c.Name.StartsWith ("T")
from p in c.Purchases
select new { c.Name, p.Description };
```

This LINQ-to-db query would work equally well if the where clause is moved one line down. If it is a local query, however, moving the where clause would make it less efficient. With local queries, you should filter *before* joining.

You can introduce new tables into the mix with additional from clauses. For instance, if each purchase had purchase item child rows, you could produce a flat result set of customers with their purchases, each with their purchase detail lines as follows:

```
from c in DataContext.Customers
from p in c.Purchases
from pi in p.PurchaseItems
select new { c.Name, p.Description, pi.DetailLine };
```

Each from clause introduces a new *child* table. To include data from a *parent* table (via an association property), you don't add a from clause—you simply navigate to

Each `from` clause introduces a new *child* table. To include data from a *parent* table (via an association property), you don't add a `from` clause—you simply navigate to the property. For example, if each customer has a salesperson whose name you want to query, just do this:

```
from c in dataContext.Customers
select new { Name = c.Name, SalesPerson = c.SalesPerson.Name };
```

You don't use `SelectMany` in this case because there's no subcollection to flatten. Parent association properties return a single item.

## Outer joins with `SelectMany`

We saw previously that a `Select` subquery yields a result analogous to a left outer join:

```
from c in dataContext.Customers
select new {
    c.Name,
```

```
Purchases = from p in c.Purchases
```

Purchases = from p in c.Purchases  
where p.Price > 1000

ΛINQ Operators

};



```
};
```

```
select new { p.Description, p.Price }
```

In this example, every outer element (customer) is included, regardless of whether the customer has any purchases. But suppose we rewrite this query with `SelectMany`, so we can obtain a single flat collection rather than a hierarchical result set:

```
from c in DataContext.Customers  
from p in c.Purchases  
where p.Price > 1000  
select new { c.Name, p.Description, p.Price };
```

In the process of flattening the query, we've switched to an inner join: customers are now included only for whom one or more high-value purchases exist. To get a left outer join with a flat result set, we must apply the `DefaultIfEmpty` query operator on the inner sequence. This method returns null if its input sequence has no elements. Here's such a query, price predicate aside:

ments. Here's such a query, price predicate aside:

```
from c in DataContext.Customers
from p in c.Purchases.DefaultIfEmpty()
select new { c.Name, p.Description, Price = (decimal?) p.Price };
```

This works perfectly with LINQ to SQL and EF, returning all customers, even if they have no purchases. But if we were to run this as a local query, it would crash, because when `p` is null, `p.Description` and `p.Price` throw a `NullReferenceException`. We can make our query robust in either scenario as follows:

```
from c in DataContext.Customers
from p in c.Purchases.DefaultIfEmpty()
select new {
    c.Name,
    Descript = p == null ? null : p.Description,
    Price = p == null ? (decimal?) null : p.Price
};
```

Let's now reintroduce the price filter. We cannot use a `where` clause as we did before, because it would execute *after* `DefaultIfEmpty`:

Let's now reintroduce the price filter. We cannot use a `where` clause as we did before, because it would execute *after* `DefaultIfEmpty`:

```
from c in DataContext.Customers
from p in c.Purchases.DefaultIfEmpty()
where p.Price > 1000...
```

The correct solution is to splice the `where` clause *before* `DefaultIfEmpty` with a subquery:

```
from c in DataContext.Customers
from p in c.Purchases.Where (p => p.Price > 1000).DefaultIfEmpty()
select new {
    c.Name,
    Descript = p == null ? null : p.Description,
    Price = p == null ? (decimal?) null : p.Price
};
```

LINQ to SQL and EF translate this to a left outer join. This is an effective pattern for writing such queries.

LINQ to SQL and LINQ to Entities and to write outer joins. LINQ is an effective pattern for writing such queries.



If you're used to writing outer joins in SQL, you might be tempted to overlook the simpler option of a `Select` subquery for this style of query, in favor of the awkward but familiar SQL-centric flat approach. The hierarchical result set from a `Select` subquery is often better suited to outer join-style queries because there are no additional nulls to deal with.

# Joining

`IEnumerable<TOuter>`, `IEnumerable<TInner>` → `IEnumerable<TResult>`

Method

Description

SQL equivalents

Join

Applies a lookup strategy to match elements from two collections, emitting

INNER JOIN

Join	Applies a lookup strategy to match elements from two collections, emitting a flat result set	INNER JOIN
------	--	------------

## GroupJoin

As above, but emits a *hierarchical* result set

INNER JOIN,

LEFT OUTER JOIN

# Join and GroupJoin

## Join arguments

Argument	Type
----------	------

Outer sequence  
IF [enumerable] < TOuter >

Outer sequence	IEnumerable<TOuter>
Inner sequence	IEnumerable<TInner>
Outer key selector	TOuter => TKey
Inner key selector	TInner => TKey
Result selector	(TOuter, TInner) => TResult

# GroupJoin arguments

Argument	Type
Outer sequence	IEnumerable<TOuter>
Inner sequence	IEnumerable<TInner>
Outer key selector	TOuter => TKey
Inner key selector	TInner => TKey

Inner key selector    TInner => TKey

Result selector    (TOuter, IEnumerable<TInner>) => TResult

---

# Query syntax

*from outer-var in outer-enumerable*  
*join inner-var in inner-enumerable on outer-key-expr equals inner-key-expr*  
*[ into identifier ]*

## Overview

`Join` and `GroupJoin` mesh two input sequences into a single output sequence. `Join` emits flat output; `GroupJoin` emits hierarchical output.

`Join` and `GroupJoin` provide an alternative strategy to `Select` and `SelectMany`. The advantage of `Join` and `GroupJoin` is that they execute efficiently over local in-memory collections, since they first load the inner sequence into a keyed lookup, avoiding the need to repeatedly enumerate over every inner element. The disadvantage is that they offer the equivalent of inner and left outer joins only; cross joins and non-equijoin joins must still be done with `Select/SelectMany`. With LINQ to SQL and Entity Framework queries, `Join` and `GroupJoin` offer no real benefits over `Select` and `SelectMany`.

Table 9-1 summarizes the differences between each of the joining strategies.



Table 9-1 summarizes the differences between each of the joining strategies.

Table 9-1. *Joining strategies*

Strategy	Left					Non-
	Result shape	Local query efficiency	Inner joins	outer joins	Cross joins	equi joins
Select + SelectMany	Flat	Bad	Yes	Yes	Yes	Yes
Select + Select	Nested	Bad	Yes	Yes	Yes	Yes

Join

Flat

Good

Yes

-  
Yes

Yes

-  
-  
-

GroupJoin

Nested

Good

Yes

GroupJoin + SelectMany

Flat

Good

Yes

Yes

-

-

Join

The Join operator performs an inner join, emitting a flat output sequence.

The Join operator performs an inner join, emitting a flat output sequence.



Entity Framework hides foreign key fields, so you can't manually join across natural relationships (instead, you can query across association properties, as we described in the previous two sections).

The simplest way to demonstrate Join is with LINQ to SQL. The following query lists all customers alongside their purchases, without using an association property:

```
Queryable<string> query =  
    from c in DataContext.Customers  
    join p in DataContext.Purchases on c.ID equals p.CustomerID  
    select c.Name + " bought a " + p.Description;
```

The results match what we would get from a SelectMany-style query:

Tom bought a Bike

Tom bought a Phone  
Tom bought a Holiday  
Dick bought a Phone  
Harry bought a Car

---

## 388 | Chapter 9: LINQ Operators

---

To see the benefit of Join over SelectMany, we must convert this to a local query. We can demonstrate this by first copying all customers and purchases to arrays, and then querying the arrays:

```
Customer[] customers = dataContext.Customers.ToArray();  
Purchase[] purchases = dataContext.Purchases.ToArray();  
  
var slowQuery = from c in customers  
                from p in purchases where c.ID == p.CustomerID  
                select c.Name + " bought a " + p.Description;
```

```
var fastQuery = from c in customers  
                join p in purchases on c.ID equals p.CustomerID  
                select c.Name + " bought a " + p.Description;
```

```
select c.Name + " bought a " + p.Description;
```

Although both queries yield the same results, the Join query is considerably faster because its implementation in Enumerable preloads the inner collection (*purchases*) into a keyed lookup.

The query syntax for join can be written in general terms as follows:

join *inner-var* in *inner-sequence* on *outer-key-expr* equals *inner-key-expr*

Join operators in LINQ differentiate between the *outer sequence* and *inner sequence*. Syntactically:

- 
- 

The *outer sequence* is the input sequence (in this case, *customers*).

The *inner sequence* is the new collection you introduce (in this case, *purchases*).

Join performs inner joins, meaning customers without purchases are excluded from the output. With inner joins, you can swap the inner and outer sequences in the query and still get the same results:

the output. With inner joins, you can swap the inner and outer sequences in the query and still get the same results:

```
from p in purchases                                // p is now outer
join c in customers on p.CustomerID equals c.ID    // c is now inner
...
```

You can add further join clauses to the same query. If each purchase, for instance, has one or more purchase items, you could join the purchase items as follows:

```
from c in customers
join p in purchases on c.ID equals p.CustomerID
join pi in purchaseItems on p.ID equals pi.PurchaseID
...
```

```
// first join
// second join
```

`purchases` acts as the *inner* sequence in the first join and as the *outer* sequence in the second join. You could obtain the same results (inefficiently) using nested `foreach` statements as follows:

second join. You could obtain the same results (inefficiency) using nested foreach statements as follows:

```
foreach (Customer c in customers)
    foreach (Purchase p in purchases)
        if (c.ID == p.CustomerID)
            foreach (PurchaseItem pi in purchaseItems)
                if (p.ID == pi.PurchaseID)
                    Console.WriteLine (c.Name + ", " + p.Price + ", " + pi.Detail);
```

In query syntax, variables from earlier joins remain in scope—just as outer range variables do with `SelectMany`-style queries. You're also permitted to insert `where` and `let` clauses in between join clauses.

## Joining on multiple keys

You can join on multiple keys with anonymous types as follows:

```
from x in sequenceX
join y in sequenceY on new { K1 = x.Prop1, K2 = x.Prop2 }
equals new { K1 = y.Prop3, K2 = y.Prop4 }
...
```

For this to work, the two anonymous types must be structured identically. The compiler then implements each with the same internal type, making the joining keys compatible.



compiler then implements each with the same internal type, making the joining keys compatible.

## Joining in fluent syntax

The following query syntax join:

```
from c in customers
join p in purchases on c.ID equals p.CustomerID
select new { c.Name, p.Description, p.Price };
```

in fluent syntax is as follows:

```
customers.Join (
    purchases,
    c => c.ID,
    p => p.CustomerID,
    (c, p) => new
        { c.Name, p.Description, p.Price } // result selector
);
```

```
    { c.Name, p.Description, p.Price } // result selector  
  );
```

The result selector expression at the end creates each element in the output sequence. If you have additional clauses prior to projecting, such as **orderby** in this example:

```
from c in customers  
join p in purchases on c.ID equals p.CustomerID  
orderby p.Price  
select c.Name + " bought a " + p.Description;
```

you must manufacture a temporary anonymous type in the result selector in fluent syntax. This keeps both *c* and *p* in scope following the join:

```
customers.Join (           // outer collection  
    purchases,           // inner collection  
    c => c.ID,           // outer key selector  
    p => p.CustomerID,   // inner key selector  
    (c, p) => new { c, p } ) // result selector  
    .Select (x => x.c.Name + " bought a " + x.p.Description);
```

Query syntax is usually preferable when joining; it's less fiddly.

## GroupJoin

GroupJoin does the same work as Join, but instead of yielding a flat result, it yields a hierarchical result, grouped by each outer element. It also allows left outer joins.

The query syntax for GroupJoin is the same as for Join, but is followed by the `into` keyword.

Here's the most basic example:

```
IEnumerable<IEnumerable<Purchase>> query =  
    from c in customers  
    join p in purchases on c.ID equals p.CustomerID  
    into custPurchases
```

into **custPurchases**

```
select custPurchases; // custPurchases is a sequence
```



An **into** clause translates to **GroupJoin** only when it appears directly after a **join** clause. After a **select** or **group** clause, it means *query continuation*. The two uses of the **into** keyword are quite different, although they have one feature in common: they both introduce a new query variable.

The result is a sequence of sequences, which we could enumerate as follows:

```
foreach (IEnumerable<Purchase> purchasesSequence in query)
    foreach (Purchase p in purchasesSequence)
        Console.WriteLine (p.Description);
```

This isn't very useful, however, because **outerSeq** has no reference to the outer customer. More commonly, you'd reference the outer range variable in the projection:

```
from c in customers
join p in purchases on c.ID equals p.CustomerID
```

```
join p in purchases on c.ID equals p.CustomerID  
into custPurchases  
select new { CustName = c.Name, custPurchases };
```

This gives the same results as the following (inefficient) Select subquery:

```
from c in customers  
select new  
{  
    CustName = c.Name,  
    custPurchases = purchases.Where (p => c.ID == p.CustomerID)  
};
```

By default, GroupJoin does the equivalent of a left outer join. To get an inner join—where customers without purchases are excluded—filter on custPurchases:

```
from c in customers join p in purchases on c.ID equals p.CustomerID  
into custPurchases  
where custPurchases.Any()  
select ...
```

```
WHERE cust IN (SELECT c.custid FROM cust c)  
select ...
```

Clauses after a group-join into operate on *subsequences* of inner child elements, not *individual* child elements. This means that to filter individual purchases, you'd have to call *where before* joining:

## JOIN Operators

```
from c in customers
join p in purchases where (p2 => p2.Price > 1000)
on c.ID equals p.CustomerID
into custPurchases ...
```

You can construct lambda queries with GroupJoin as you would with Join.

## Flat outer joins

You run into a dilemma if you want both an outer join and a flat result set. Group Join gives you the outer join; Join gives you the flat result set. The solution is to first call GroupJoin, and then DefaultIfEmpty on each child sequence, and then finally SelectMany on the result:

```
from c in customers
join p in purchases on c.ID equals p.CustomerID into custPurchases
from cp in custPurchases.DefaultIfEmpty()
select new
```

```
select new  
{  
    CustName = c.Name,  
    Price = cp == null ? (decimal?) null : cp.Price  
};
```

DefaultIfEmpty emits a null value if a subsequence of purchases is empty. The second from clause translates to `SelectMany`. In this role, it *expands and flattens* all the purchase subsequences, concatenating them into a single sequence of purchase *elements*.

## Joining with lookups

The `Join` and `GroupJoin` methods in `Enumerable` work in two steps. First, they load the inner sequence into a *lookup*. Second, they query the outer sequence in combination with the lookup.

A *lookup* is a sequence of groupings that can be accessed directly by key. Another way to think of it is as a dictionary of sequences—a dictionary that can accept many elements under each key. Lookups are read-only and defined by the following interface:



elements under each key. Lookups are read-only and defined by the following interface:

```
public interface ILookup<TKey, TElement> :  
    IEnumerable<IGrouping<TKey, TElement>>, IEnumerable  
{  
    int Count { get; }  
    bool Contains (TKey key);  
    IEnumerable<TElement> this [TKey key] { get; }  
}
```



The joining operators—like other sequence-emitting operators—honor deferred or lazy execution semantics. This means the lookup is not built until you begin enumerating the output sequence.

You can create and query lookups manually as an alternative strategy to using the joining operators, when dealing with local collections. There are a couple of benefits in doing so:

- 
- 

You can reuse the same lookup over multiple queries—as well as in ordinary imperative code.

Querying a lookup is an excellent way of understanding how `Join` and `GroupJoin` work.

The `ToLookup` extension method creates a lookup. The following loads all purchases into a lookup—keyed by their `CustomerID`:

```
ILookup<int?, Purchase> purchLookup =  
    purchases.ToLookup (p => p.CustomerID, p => p);
```

The first argument selects the key; the second argument selects the objects that are to be loaded as values into the lookup.

to be loaded as values into the lookup.

Reading a lookup is rather like reading a dictionary, except that the indexer returns a *sequence* of matching items, rather than a *single* matching item. The following enumerates all purchases made by the customer whose ID is 1:

```
foreach (Purchase p in purchlookup [1])  
    Console.WriteLine (p.Description);
```

With a lookup in place, you can write `SelectMany/Select` queries that execute as efficiently as `Join/GroupJoin` queries. `Join` is equivalent to using `SelectMany` on a lookup:

```
from c in customers  
from p in purchlookup [c.ID]  
select new { c.Name, p.Description, p.Price };
```

Tom Bike 500

Tom Holiday 2000

Dick Bike 600

Dick Bike 600  
Dick Phone 300  
...

Adding a call to `DefaultIfEmpty` makes this into an outer join:

```
from c in customers
from p in purchlookup [c.ID].DefaultIfEmpty()
select new {
    c.Name,
    Descript = p == null ? null : p.Description,
    Price = p == null ? (decimal?) null : p.Price
};
```

# NQ Operators

GroupJoin is equivalent to reading the lookup inside a projection:

```
from c in customers
select new {
    CustName = c.Name,
    CustPurchases = purchLookup [c.ID]
};
```

Enumerable implementations

## Enumerable implementations

Here's the simplest valid implementation of Enumerable.Join, null checking aside:

```
public static IEnumerable<TResult> Join
    <TOuter, TInner, TKey, TResult> (
    this IEnumerable<TOuter>    outer,
    IEnumerable<TInner>        inner,
    Func<TOuter, TKey>          outerKeySelector,
    Func<TInner, TKey>          innerKeySelector,
    Func<TOuter, TInner, TResult> resultSelector)
{
    Ilookup<TKey, TInner> lookup = inner.ToLookup (innerKeySelector);
    return
        from outerItem in outer
        from innerItem in lookup [outerKeySelector (outerItem)]
        select resultSelector (outerItem, innerItem);
}
```

GroupJoin's implementation is like that of Join, but simpler:

```
public static IEnumerable<TResult> GroupJoin
```

```

public static IEnumerable <TResult> GroupJoin
    <TOuter, TInner, TKey, TResult> (
        this IEnumerable <TOuter>      outer,
        IEnumerable <TInner>           inner,
        Func <TOuter, TKey>             outerKeySelector,
        Func <TInner, TKey>             innerKeySelector,
        Func <TOuter, IEnumerable<TInner>, TResult> resultSelector)
{
    Ilookup <TKey, TInner> lookup = inner.ToLookup (innerKeySelector);
    return
        from outerItem in outer
        select resultSelector
            (outerItem, lookup [outerKeySelector (outerItem)]);
}

```

# Ordering

`IEnumerable<TSource>→IOOrderedEnumerable<TSource>`

Method

Description

SQL equivalents

Method	Description	SQL equivalents
OrderBy, ThenBy	Sorts a sequence in ascending order	ORDER BY ...
OrderByDescending, ThenByDescending	Sorts a sequence in descending order	ORDER BY ... DESC
Reverse	Returns a sequence in reverse order	Exception thrown

Ordering operators return the same elements in a different order.

## OrderBy, OrderByDescending, ThenBy, and ThenByDescending

### OrderBy and OrderByDescending arguments

Argument	Type
Input sequence	<code>IEnumerable&lt;TSource&gt;</code>
Key selector	<code>TSource =&gt; TKey</code>

Return type – `IOderedEnumerable<TSource>`



Return type = `IOrderedEnumerable<TSource>`

## ThenBy and ThenByDescending arguments

Argument	Type
Input sequence	<code>IOrderedEnumerable&lt;TSource&gt;</code>
Key selector	<code>TSource =&gt; TKey</code>

### Query syntax

```
orderby expression1 [descending] [, expression2 [descending] ... ]
```

## Overview

`OrderBy` returns a sorted version of the input sequence, using the `keySelector` expression to make comparisons. The following query emits a sequence of names in alphabetical order:

alphabetical order:

```
IEnumerable<string> query = names.OrderBy (s => s);
```

The following sorts names by length:

```
IEnumerable<string> query = names.OrderBy (s => s.Length);
```

```
// Result: { "Jay", "Tom", "Mary", "Dick", "Harry" };
```

The relative order of elements with the same sorting key (in this case, Jay/Tom and Mary/Dick) is indeterminate—unless you append a `ThenBy` operator:

```
IEnumerable<string> query = names.OrderBy (s => s.Length).ThenBy (s => s);
```

```
// Result: { "Jay", "Tom", "Dick", "Mary", "Harry" };
```

`ThenBy` reorders only elements that had the same sorting key in the preceding sort. You can chain any number of `ThenBy` operators. The following sorts first by length, then by the second character, and finally by the first character:

```
names.OrderBy (s => s.Length).ThenBy (s => s[1]).ThenBy (s => s[0]);
```

The `OrderBy` method in `System.Linq` is this:

The equivalent in query syntax is this:

```
from s in names  
orderby s.Length, s[1], s[0]  
select s;
```

LINQ also provides `OrderByDescending` and `ThenByDescending` operators, which do the same things, emitting the results in reverse order. The following LINQ-to-db query retrieves purchases in descending order of price, with those of the same price listed alphabetically:

```
dataContext.Purchases.OrderByDescending (p => p.Price)  
    .ThenBy (p => p.Description);
```

## In query syntax:

```
from p in dataContext.Purchases  
orderby p.Price descending, p.Description  
select p;
```

### Comparers and collations

In a local query, the key selector objects themselves determine the ordering algorithm

**comparer** and **insensitive** attributes.

In a local query, the key selector objects themselves determine the ordering algorithm via their default `IComparable` implementation (see Chapter 7). You can override the sorting algorithm by passing in an `IComparer` object. The following performs a case-insensitive sort:

```
names.OrderBy (n => n, StringComparer.CurrentCultureIgnoreCase);
```

Passing in a comparer is not supported in query syntax, nor in any way by LINQ to SQL or EF. When querying a database, the comparison algorithm is determined by the participating column's collation. If the collation is case-sensitive, you can request a case-insensitive sort by calling `Toupper` in the key selector:

```
from p in DataContext.Purchases  
orderby p.Description.Toupper()  
select p;
```

# **IOrderedEnumerable and IOrderedQueryable**

The ordering operators return special subtypes of `IEnumerable<T>`. Those in `Enumerable` return `IOrderedEnumerable`; those in `Queryable` return `IOrderedQueryable`. These subtypes allow a subsequent `ThenBy` operator to refine rather than replace the existing ordering.

The additional members that these subtypes define are not publicly exposed, so they present like ordinary sequences. The fact that they are different types comes into play when building queries progressively:

```
IOrderedEnumerable<string> query1 = names.OrderBy (s => s.Length);  
IOrderedEnumerable<string> query2 = query1.ThenBy (s => s);
```

If we instead declare `query1` of type `IEnumerable<string>`, the second line would not compile—`ThenBy` requires an input of type `IOrderedEnumerable<string>`. You can avoid worrying about this by implicitly typing query variables:

```
var query1 = names.OrderBy (s => s.Length);  
var query2 = query1.ThenBy (s => s);
```

Implicit typing can create problems of its own, though. The following will not compile:

```
var query = names.OrderBy (s => s.Length);  
query = query.Where (n => n.Length > 3);
```

## // Compile-time error

The compiler infers query to be of type `IOrderedEnumerable<string>`, based on `OrderBy`'s output sequence type. However, the `Where` on the next line returns an ordinary `IEnumerable<string>`, which cannot be assigned back to query. You can work around this either with explicit typing or by calling `AsEnumerable()` after `OrderBy`:

```
var query = names.OrderBy (s => s.Length).AsEnumerable();  
query = query.Where (n => n.Length > 3);           // OK
```

The equivalent in interpreted queries is to call `AsQueryable`.

# Grinning

# Grouping

`IEnumerable<TSource> → IEnumerable<IGrouping<TSource, TElement>>`

Method	Description	SQL equivalents
GroupBy	Groups a sequence into subsequences	GROUP BY

## GroupBy

Argument	Type
Input sequence	<code>IEnumerable&lt;TSource&gt;</code>
Key selector	<code>TSource =&gt; TKey</code>
Element selector (optional)	<code>TSource =&gt; TElement</code>
Comparer (optional)	<code>IEqualityComparer&lt;TKey&gt;</code>



## Query syntax

`group element-expression by key-expression`

## Overview

GroupBy organizes a flat input sequence into sequences of *groups*. For example, the following organizes all the files in *c:\temp* by extension:

```
string[] files = Directory.GetFiles ("c:\\temp");
```

```
IEnumerable<IGrouping<string,string>> query =  
files.GroupBy (file => Path.GetExtension (file));
```

Or if you're comfortable with implicit typing:

```
var query = files.GroupBy (file => Path.GetExtension (file));
```

Here's how to enumerate the result:

## LINQ Operators

```
foreach (IGrouping<string,string> grouping in query)
{
    Console.WriteLine ("Extension: " + grouping.Key);
    foreach (string filename in grouping)
    {
        Console.WriteLine ("    {0} ({1})", filename, grouping.Key);
    }
}
```

```
foreach (string filename in grouping)
    Console.WriteLine ("  - " + filename);
}
```

```
Extension: .pdf
```

```
-- chapter03.pdf
```

```
-- chapter04.pdf
```

```
Extension: .doc
```

```
-- todo.doc
```

```
-- menu.doc
```

```
-- Copy of menu.doc
```

```
...
```

---

`Enumerable.GroupBy` works by reading the input elements into a temporary dictionary of lists so that all elements with the same key end up in the same sublist. It then emits a sequence of *groupings*. A grouping is a sequence with a `Key` property:

---

```
public interface IGrouping<TKey,TElement> : IEnumerable<TElement>,
    IEnumerable
```

```
{
```

```
{  
    TKey Key { get; }    // Key applies to the subsequence as a whole  
}
```

By default, the elements in each grouping are untransformed input elements, unless you specify an `elementSelector` argument. The following projects each input element to uppercase:

```
files.GroupBy (file => Path.GetExtension (file), file => file.ToUpper());
```

An `elementSelector` is independent of the `keySelector`. In our case, this means that the `Key` on each grouping is still in its original case:

```
Extension: .pdf  
-- CHAPTER03.PDF  
-- CHAPTER04.PDF  
Extension: .doc  
-- TODO.DOC
```

Note that the subcollections are not emitted in alphabetical order of `key`. `GroupBy` groups only; it does not *sort*; in fact, it preserves the original ordering. To sort, you must add an `OrderBy` operator:

```
files.GroupBy (file => Path.GetExtension (file), file => file.ToUpper())  
    .OrderBy (grouping => grouping.Key);
```

```
files.GroupBy (file => Path.GetExtension (file), file => file.ToUpper())  
    .OrderBy (grouping => grouping.Key);
```

GroupBy has a simple and direct translation in query syntax:

***group element-expr by key-expr***

Here's our example in query syntax:

```
from file in files  
group file.ToUpper() by Path.GetExtension (file);
```

As with select, group “ends” a query—unless you add a query continuation clause:

---

398 | Chapter 9: LINQ Operators

```
from file in files  
group file.ToUpper() by Path.GetExtension (file) into grouping  
orderby grouping.Key
```

orderby grouping.Key  
select grouping;

Query continuations are often useful in a group by query. The next query filters out groups that have fewer than five files in them:

```
from file in files  
group file.Tolpper() by Path.GetExtension (file) into grouping  
where grouping.Count() < 5  
select grouping;
```



A where after a group by is equivalent to HAVING in SQL. It applies to each subsequence or grouping as a whole, rather than the individual elements.

Sometimes you're interested purely in the result of an aggregation on a grouping and so can abandon the subsequences:

```
string[] votes = { "Bush", "Gore", "Gore", "Bush", "Bush" };
```

```
IEnumerable<string> query = from vote in votes  
                             group vote by vote into g  
                             orderby g.Count() descending  
                             select g.Key;
```

```
string winner = query.First();
```

```
// Bush
```

## GroupBy in LINQ to SQL and EF

Grouping works in the same way when querying a database. If you have association properties set up, you'll find, however, that the need to group arises less frequently than with standard SQL. For instance, to select customers with at least two purchases, you don't need to group; the following query does the job nicely:

```
from c in DataContext.Customers
```

```
from c in dataContext.Customers
where c.Purchases.Count >= 2
select c.Name + " has made " + c.Purchases.Count + " purchases";
```

An example of when you might use grouping is to list total sales by year:

```
from p in dataContext.Purchases
group p.Price by p.Date.Year into salesByYear
select new {
    Year      = salesByYear.Key,
    TotalValue = salesByYear.Sum()
};
```

LINQ's grouping operators expose a superset of SQL's "GROUP BY" functionality.

Another departure from traditional SQL comes in there being no obligation to project the variables or expressions used in grouping or sorting.



# LINQ Operators

## Grouping by multiple keys

You can group by a composite key, using an anonymous type:

```
from n in names
group n by new { FirstLetter = n[0], Length = n.Length };
```

## Custom equality comparers

You can pass a custom equality comparer into `GroupBy`, in a local query, to change the algorithm for key comparison. Rarely is this required, though, because changing the key selector expression is usually sufficient. For instance, the following creates a case-insensitive grouping:

```
group name by name.ToUpper()
```

# Set Operators

`IEnumerable<TSource>`, `IEnumerable<TSource>` → `IEnumerable<TSource>`

Concat	Returns a concatenation of elements in each of the two sequences	UNION ALL
--------	--	-----------

## Union

## Intersect

## Except

Returns a concatenation of elements in each of the two sequences, excluding duplicates

## UNION

Returns elements present in both sequences

Returns elements present in the first, but not the second sequence

## WHERE ... IN (...)

**WHERE ... IN (...)**

**EXCEPT**

*or*

---

**WHERE ... NOT IN (...)**

# Concat and Union

Concat returns all the elements of the first sequence, followed by all the elements of the second. Union does the same, but removes any duplicates:

```
int[] seq1 = { 1, 2, 3 }, seq2 = { 3, 4, 5 };
```

**IEnumerable<int>**

```
concat = seq1.Concat (seq2),
```

```
union  = seq1.Union (seq2);
```